

# **Vectorization, Normalization, Outliers Detection**

**Data Analysis with Python**

**Burak Gündüz**

[github.com/burakgunduztr](https://github.com/burakgunduztr)

# **EDA (Exploratory Data Analysis)**

We're in the era of fast code and data flows, where understanding the quality of that stream is much more important than ever.

**Data Analysis with Python**

# Python list

**1,000,000 pointer (any object)**

reference count (8 bytes)  
pointer to its type (8 bytes)  
the actual 8-byte double-precision float  
plus padding...

24-32 bytes each, around 36 MB

```
# Python list - can mix types
py_list = [1, 'hello', 3.14, None]

# NumPy array - single type
import numpy as np
np_array = np.array([1, 2, 3, 4]) # all int64 (or float64)

# Common patterns
np.array([1, 2, 3], dtype=np.int32)      # 32-bit integers
np.array([1, 2, 3], dtype=np.float64)     # 64-bit floats (default)
np.array([1, 2, 3], dtype='float32')       # String notation
np.array([1, 2, 3]).astype(np.float32)     # Convert existing array

# Check type
arr.dtype

# Convert type
arr.astype(np.float64)
```

# NumPy array

Large library of array functions, fast compact (fixed type)

**import numpy as np**

**1,000,000 floats (float64)**

8 bytes each, around 8 MB

**Popular use-cases:**

*EDA, Science, Statistic*

*Linear algebra*

*FFT (Fast Fourier Transform)*

*Random generations*

# Vectorized Operations

Expressing operations to whole array in C-level, instead Python loops

```
# NON-VECTORIZED (Python list with loop)
start = time.time()
squared_list = []
for x in py_list:
    squared_list.append(x ** 2)
python_time = time.time() - start
print(f"Python loop: {python_time:.4f} seconds")

# NON-VECTORIZED (Python list comprehension – still a loop!)
start = time.time()
squared_list = [x ** 2 for x in py_list]
list_comp_time = time.time() - start
print(f"List comprehension: {list_comp_time:.4f} seconds")

# VECTORIZED (NumPy)
start = time.time()
squared_array = np_array ** 2
numpy_time = time.time() - start
print(f"NumPy vectorized: {numpy_time:.4f} seconds")

print(f"\nSpeedup: {python_time / numpy_time:.1f}x faster")
```

```
import numpy as np
import time


# Setup data
size = 1_000_000
py_list = list(range(size))
np_array = np.arange(size)
```

*np.arange(size) creates integer array including given number of elements*


Python loop: 0.1234 seconds  
List comprehension: 0.0856 seconds  
NumPy vectorized: 0.0012 seconds  
Speedup: 102.8x faster

*NumPy calculates internally in C tight row buffer (or SIMD) and then assigning result at the end*

Traditional loop (one at a time):



SIMD (parallel processing):



Data Analysis with Python

# Quick statistic checkup

Statistic	Formula
Mean ( $\mu$ )	$\mu = \frac{1}{N} \sum_{i=1}^N x_i$
Variance ( $\sigma^2$ )	$\sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$
Standard Deviation ( $\sigma$ )	$\sigma = \sqrt{\sigma^2} = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}$

*Mean:* Average of all values.

Represents the center point of the data.

*Variance:* Average squared deviation from mean.

Measures how spread out the data is. Units are squared (e.g.,  $cm^2$ ). Always  $\geq 0$

*Standard Deviation:* Square root of variance. Measures spread in original units (e.g.,  $cm$ ).

More interpretable than variance.

```
data = np.array([2, 4, 4, 4, 5, 5, 7, 9])
```

```
# Mean
```

```
mean = np.mean(data)  
print(f"Mean: {mean}") # → 5.0
```

```
# Variance
```

```
variance = np.var(data)  
print(f"Variance: {variance}") # → 4.0
```

```
# Standard Deviation
```

```
std = np.std(data)  
print(f"Std Dev: {std}") # → 2.0
```

# Why Normalization matters?

*Same scale for different features*

*having “mean 0, std 1” over complete array  
makes them comparable, none of them dominance*

*Easier for many ML algorithms*

*k-NN (k-Nearest Neighbors)  
classification and regression that makes predictions  
PCA (Principal Component Analysis)  
transforms high-dimensional data into a lower-dimensional  
helps optimization and convergence*

*Interpretation*

*Normalized 2.0 “2 std above mean”  
0.5 “half a standard deviation below”  
how unusual, how far from average*

# Pandas and its DataFrame

One object (container) holds data in a table

```
import pandas as pd
```

Operation	axis=0 (rows)	axis=1 (columns)
Direction	Down ↓	Across →
Sum	Sum each column	Sum each row
Mean	Mean of each column	Mean of each row
Delete	Delete a row	Delete a column
Concatenate	Stack vertically (more rows)	Stack horizontally (more columns)

```
import pandas as pd

# Sales data
sales = pd.DataFrame({
    'Jan': [100, 150, 200],
    'Feb': [110, 160, 210],
    'Mar': [120, 170, 220]
}, index=['Product A', 'Product B', 'Product C'])

print(sales)
#           Jan  Feb  Mar
# Product A 100  110  120
# Product B 150  160  170
# Product C 200  210  220

# axis=0: Total sales per month (sum DOWN each column)
monthly_totals = sales.sum(axis=0)
print(monthly_totals)
# Jan      450
# Feb      480
# Mar      510

# axis=1: Total sales per product (sum ACROSS each row)
product_totals = sales.sum(axis=1)
print(product_totals)
# Product A    330
# Product B    480
# Product C    630
```

# .groupBy() and .agg()

Code:

```
df.groupby('Product').agg(  
    total_sales=('Sales', 'sum'),  
    avg_quantity=('Quantity', 'mean'),  
    min_cost=('Cost', 'min'),  
    max_cost=('Cost', 'max'))
```

Output:

Product	total_sales	avg_quantity	min_cost	max_cost
A	420	2.67	60	120
B	490	3.33	90	110
C	200	1.50	50	65

Code:

```
df.groupby('Product')['Sales'].agg([  
    'sum', 'mean', 'count', 'min', 'max', 'std', 'median'])
```

Output:

Product	sum	mean	count	min	max	std	median
A	420	140.0	3	100	200	52.92	120.0
B	490	163.33	3	150	180	15.28	160.0
C	200	100.0	2	90	110	14.14	100.0

# .groupBy() and .agg()

Code:

```
df.groupby(['Product', 'Region']).agg({  
    'Sales': ['sum', 'mean'],  
    'Quantity': 'sum'  
})
```

Output:

Product	Region	Sales sum	Sales mean	Quantity sum
A	North	220	110.0	4
A	South	200	200.0	4
B	North	310	155.0	7
B	South	180	180.0	3
C	North	110	110.0	2
C	South	90	90.0	1

Code:

```
# Filter BEFORE groupby  
df[df['Sales'] > 100].groupby('Product')['Sales'].sum()
```

Output:

Product	Sales
A	320
B	490
C	110

Goal	Code
Single aggregation	df.groupby('col')['value'].sum()
Multiple aggregations	df.groupby('col')['value'].agg(['sum', 'mean'])
Different per column	df.groupby('col').agg({'col1': 'sum', 'col2': 'mean'})
Named aggregations	df.groupby('col').agg(total=('value', 'sum'))
Multiple groups	df.groupby(['col1', 'col2'])['value'].sum()
All statistics	df.groupby('col')['value'].describe()
Pivot view	df.pivot_table(values='val', index='row', columns='col')

# Why data cleaning matters?

*Sources are messy*

*Bad data breaks analysis*

*aggregations, models, visualizations*

*Supporting system expectations*

*reporting, search, recommendation engines*

# Data cleaning pipeline

*Handle missing values*

*Remove duplicates*

*Flag and process unexpected values*

*Add a new column If needed*

```
import pandas as pd

# Original data
data = {
    'Name': ['Alice', 'Bob', '', 'David'],
    'Age': [25, '', 35, 28],
    'Salary': ['50000', '60000', '', '55000'],
    'Dept': ['Sales', 'IT', 'N/A', 'Sales']
}
df = pd.DataFrame(data)

# Step 1: Replace placeholders with NA
df = df.replace(['', 'N/A'], pd.NA)

# Step 2: Fill missing values
df = df.fillna({'Name': 'Unknown', 'Age': '30', 'Salary': '0', 'Dept': 'Other'})

# Step 3: Convert types
df['Age'] = df['Age'].astype(int)
df['Salary'] = df['Salary'].astype(int)
```

# Data cleaning pipeline

## Approach 1: `fillna()` + `astype()` (What I Used)

```
# Works if you KNOW the data format is valid after filling
df = df.replace(['', 'N/A'], pd.NA)
df = df.fillna({'Age': '30', 'Salary': '0'}) # Fill with string numbers
df['Age'] = df['Age'].astype(int) # Safe – all values are numeric strings
df['Salary'] = df['Salary'].astype(int)
```

**When to use:** Data is clean, just has missing values

*astype() has no errors= param.  
use pd.to\_numeric()*

*errors='ignore' #leave as-is  
'coerce' #NaN  
'raise' #Error!*

## Approach 2: `pd.to_numeric(errors='coerce')` + `fillna()` + `astype()`

```
# Better if data might have INVALID non-numeric values
df = pd.DataFrame({
    'Age': ['25', 'N/A', '35', 'invalid', '28']
})
```

```
# Step 1: Convert, coerce invalid to NaN
df['Age'] = pd.to_numeric(df['Age'], errors='coerce')
# Result: [25.0, NaN, 35.0, NaN, 28.0]
```

```
# Step 2: Fill NaN
df['Age'] = df['Age'].fillna(30)
# Result: [25.0, 30.0, 35.0, 30.0, 28.0]
```

```
# Step 3: Convert to int (now safe!)
df['Age'] = df['Age'].astype(int)
# Result: [25, 30, 35, 30, 28]
```

**When to use:** Data might contain invalid values like 'invalid', '???' , 'unknown'

# Data cleaning pipeline

Method	Purpose	Example
replace()	Replace specific values	df.replace('N/A', np.nan)
dropna()	Remove rows with NaN	df.dropna()
dropna(subset=[...])	Remove rows with NaN in specific columns	df.dropna(subset=['Age'])
dropna(how='all')	Remove only completely empty rows	df.dropna(how='all')
fillna()	Fill NaN with value	df.fillna(0)
fillna({...})	Fill different columns differently	df.fillna({'Age': 30, 'Salary': 50000})
ffill()	Forward fill	df['col'].ffill()
bfill()	Backward fill	df['col'].bfill()
astype()	Convert data type	df['col'].astype(int)

# .apply() vs .cut()

```
# Method 1: apply() (SLOW)
def salary_bracket(salary):
    if salary < 50000:
        return 'Entry'
    elif salary < 75000:
        return 'Mid'
    elif salary < 100000:
        return 'Senior'
    else:
        return 'Executive'

start = time.time()
df['Level_Apply'] = df['Salary'].apply(salary_bracket)
time_apply = time.time() - start

# Method 2: cut() (FAST)
bins = [0, 50000, 75000, 100000, 200000]
labels = ['Entry', 'Mid', 'Senior', 'Executive']

start = time.time()
df['Level_Cut'] = pd.cut(df['Salary'], bins=bins, labels=labels)
time_cut = time.time() - start
```

## Results:

	Employee	Salary	Level_Apply	Level_Cut
0	E0000	78839	Senior	Senior
1	E0001	144670	Executive	Executive
2	E0002	52889		Mid
3	E0003	132073	Executive	Executive
4	E0004	116619	Executive	Executive

apply(): 0.0267s

cut(): 0.0008s

Speedup: 33.4x

Aspect	.apply()	.cut()
Vectorized	✗ No (row-by-row)	✓ Yes
Speed	Slow	Fast (10-50x faster)
Use case	Complex logic	Binning continuous values
Readability	Less clear	Very clear
Performance	Poor on large data	Excellent

# Quick recap for terminology

*Vectorization*

*is a way of implementing operations*

*Normalization*

*is a type of transformation*

*They're not the same process!*

*You can normalize in a vectorized way*

*or*

*Implement other (non-normalization) operations  
in a vectorized way*

# IQR (Interquartile Range)

It's a measure of spread, values are outside treated as Outliers

Finding quartiles > IQR ( $Q_3 - Q_1$ ) > Defining Fences (J.Tukey's  $1.5 \times \text{IQR}$  rule) > Filtering data gives us width of the middle 50% , spread of the mid. of the distribution

```
import numpy as np

data = np.array([2, 4, 4, 4, 5, 5, 7, 9, 15, 100])

Q1 = np.percentile(data, 25)    # 4.0
Q3 = np.percentile(data, 75)    # 8.5
IQR = Q3 - Q1                 # 4.5


lower = Q1 - 1.5 * IQR         # 4.0 - 6.75 = -2.75
upper = Q3 + 1.5 * IQR         # 8.5 + 6.75 = 15.25

outliers = data[(data < lower) | (data > upper)]
```

```
Q1      = 4.0
Q3      = 8.5
IQR     = 4.5
Lower   = -2.75
Upper   = 15.25
Outliers: [100]
```

Library	Code
NumPy	np.percentile(data, 75) - np.percentile(data, 25)
SciPy	stats.iqr(data)
Pandas	s.quantile(0.75) - s.quantile(0.25)

Recommended: `stats.iqr(data)` — cleanest, one-liner, purpose-built.



# Std

(Standart Deviation)

**measure of spread**

“typical distance” from the mean  
one number, doesn’t change each x

think a song dataset with  
mean=120 BPM, std=10 (It’s also BPM)  
allow us to call: relative tight, much wider or above

$$\sigma = \sqrt{\frac{N}{\sum_{i=1}^N (x_i - \bar{x})^2}}$$

**lower std:** values cluster near mean  
**higher std:** values spread out or multimodel

# Normalization

Z-Score Standardization (StandardScaler)

**rescale each value**

“How many standard deviations”  
above or below from the mean

$$z = \frac{x_i - \bar{x}}{\sigma}$$


(gives us std unit result)

think a tempo (BPM) dataset with  
mean=120, std=20

think a duration (sec.) dataset with  
mean=180, std=10

**allows us “scale-free comparison”**

$z = 0$	Exactly at the mean
$z > 0$	Above the mean
$z < 0$	Below the mean
$z = +1$	1 std above the mean
$z = -2$	2 stds below the mean



$abs(z) > 2$  is commonly used  
as flagged outlier

*Z-score is calculated per row after mean and std calculated in the column*

<b>Term</b>	<b>In Statistics</b>	<b>In a DataFrame</b>
<b>Column</b>	A variable / feature	<code>df['Age']</code>
<b>Row</b>	One observation / data point	A single person, transaction, etc.
<b>Mean &amp; Std</b>	Computed per <b>column</b>	One value summarizing the whole feature
<b>Z-score</b>	Computed per <b>row</b>	One score per individual observation

> Each column gets its **own** mean and std. You never mix columns when computing z-scores — Age's mean is never used to normalize Salary.


...

## Data Analysis with Python

# Normalization (Min-Max)

## Min-Max Normalization (MinMaxScaler)

$$x_{norm} = \frac{x_i - x_{min}}{x_{max} - x_{min}}$$



- The **minimum** always maps to 0
- The **maximum** always maps to 1
- Everything else falls **proportionally** in between

## Min-Max vs Z-Score

	Min-Max	Z-Score
Output range	Always 0 to 1	Typically -3 to +3
Handles outliers	Poorly — outlier pulls everything toward 0	Better — outlier just gets a large z value
Preserves zero	No — 0 in original ≠ 0 in result	No
Best for	Neural networks, image pixels, bounded features	Statistics, outlier detection, comparing features
Sensitive to new data	Yes — new min/max changes all values	Yes — new mean/std changes all values

> Use **Min-Max** when you need a strict 0–1 boundary (e.g. neural network inputs, percentage scores). Use **Z-score** when you care about how unusual a value is relative to the distribution.

# Over the Math/Programming Formulas

**Mathematical Formula (i=1 to n)**

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

---

**Programming Formula (i=0 to n-1)**

$$s^2 = \frac{1}{n-1} \sum_{i=0}^{n-1} (x_i - \bar{x})^2$$

**Where:**

- **n** = number of data points
- **n-1** = denominator (Bessel's correction, ddof=1)
- **i starts at 0** and goes to **n-1** (n iterations total)
- **x\_i** = data point at index i
- **$\bar{x}$**  = mean of the data

**Data Analysis with Python**

# DDOF (Delta Degrees of Freedom)

with Bessel's correction bias

averaging with n  
denominator is too large      unbiased estimate of population variance  
(averaging over the right count (matches the count of free pieces))

What	ddof=0 (population)	ddof=1 (sample)
Variance	$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$	$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$
Std	$\sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$	$s = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$
Z-score	$z_i = \frac{x_i - \bar{x}}{\sigma}$ (uses $\sigma$ above)	$z_i = \frac{x_i - \bar{x}}{s}$ (uses $s$ above)
Min-max [0, 1]	$\frac{x_i - x_{\min}}{x_{\max} - x_{\min}}$	(no ddof)

```
import numpy as np
from scipy import stats
from sklearn.preprocessing import StandardScaler, MinMaxScaler

data = [2, 4, 4, 4, 5, 5, 7, 9]

# Variance
np.var(data, ddof=1) # 4.571

# Standard Deviation
np.std(data, ddof=1) # 2.138

# Z-Score
stats.zscore(data, ddof=1) # [-1.402, -0.468, ..., 1.870]

# Min-Max [0, 1]
MinMaxScaler().fit_transform(np.array(data).reshape(-1, 1)) # [[0.], [0.286], ..., [1.]]
```

# Applying for Apple's Swift

**import simd**

```
// Process 4 floats simultaneously in a single CPU instruction
let a = SIMD4<Float>(1.0, 2.0, 3.0, 4.0)
let b = SIMD4<Float>(5.0, 6.0, 7.0, 8.0)

let sum      = a + b          // [6, 8, 10, 12]
let product = a * b          // [5, 12, 21, 32]
let dotProd = (a * b).sum()   // 70.0

// Normalize a vector to unit length
func normalize(_ v: SIMD4<Float>) -> SIMD4<Float> {
    return v / sqrt((v * v).sum())
}
```

**import Accelerate**

```
var data: [Float] = [10, 20, 30, 40, 50]

// --- Vectorized arithmetic ---
var result = [Float](repeating: 0, count: data.count)
var scalar: Float = 2.0
vDSP_vsmul(data, 1, &scalar, &result, 1, vDSP_Length(data.count))
// result = [20, 40, 60, 80, 100]

// --- Min-Max Normalization using vDSP ---
func vDSPNormalize(_ input: [Float]) -> [Float] {
    var minValue: Float = 0, maxValue: Float = 0
    vDSP_minv(input, 1, &minValue, vDSP_Length(input.count))
    vDSP_maxv(input, 1, &maxValue, vDSP_Length(input.count))

    var shifted = [Float](repeating: 0, count: input.count)
    var negMin = -minValue
    vDSP_vsadd(input, 1, &negMin, &shifted, 1, vDSP_Length(input.count))

    var range = maxValue - minValue
    var output = [Float](repeating: 0, count: input.count)
    vDSP_vsdiv(shifted, 1, &range, &output, 1, vDSP_Length(input.count))
    return output
}

let normalized = vDSPNormalize([10, 20, 30, 40, 50])
// [0.0, 0.25, 0.5, 0.75, 1.0]
```

Approach	Best For
Pure Swift	Small datasets, readability
SIMD	Fixed-size vectors, geometry, ML inference
vDSP (Accelerate)	Large arrays, audio, signal processing
NNNS / CreateML	Full ML pipelines on Apple platforms

**Thank you.**

**Burak Gündüz**  
**Data Analysis with Python**  
[github.com/burakgunduztr](https://github.com/burakgunduztr)